CS6208 Progress Report

Ask or Act for Cooperative Assistance via Inverse Planning

Team member: Zhu Yufan A0238993J *March 13, 2026*

What Has Been Done

All milestones through March 13 are complete. We implement a cooperative 9×9 two-room gridworld where a principal holds a hidden goal object and issues a short, potentially ambiguous instruction that maps to $K \in \{1, 2, 3, 4\}$ candidate goals. The principal acts according to a Boltzmann policy with rationality β and noise ε [1, 3]; the assistant observes states and principal actions but not the true goal.

The assistant maintains a Bayesian posterior $b(g)$ over candidate goals, initialised from the instruction template and updated from observed actions via an inverse-planning likelihood. At each step it compares the expected team cost of acting immediately versus asking the best available clarifying question:

$$\text{CostAsk}(q) = 1 + c_q + \mathbb{E}_{a|q}[V(s, b_{q,a})],$$
$$\text{CostAct} = 1 + V(s, b, g^*),$$

where c_q=0.5 is the question cost, $b_{q,a}$ is the posterior after answer a, and V is the expected BFS distance to the MAP goal [2, 4]. This is a one-step value-of-information (VoI) decision rule.

Seven policies were implemented and evaluated: *AskOrAct* (VoI), *NeverAsk*, *AlwaysAsk*, *InfoGainAsk*, *EasyInfoGainAsk*, *RandomAsk*, and *POMCPPlanner* (a POMDP tree-search baseline). The question bank contains three question types: ask color, ask room, and ask object type, with per-type answer noise (0.05, 0.10, 0.20).

A full evaluation sweep was conducted over $\varepsilon \in \{0, 0.05, 0.1\}$, $\beta \in \{1, 2, 4\}$, 5 replicate seeds $\times$ 20 episodes per condition (100 episodes per cell, 19,800 total), with 95% bootstrap confidence intervals. Additional experiments include: ablations over question cost, wrong-pick failure mode, and episode deadline margin; robustness sweeps over answer noise and principal β-mismatch; held-out instruction template generalization; and a scale-K stress test with $K \in \{5, 6\}$.

Bug discovered and fixed. During evaluation, *AskOrAct* produced results identical to *NeverAsk* at K=2, asking zero questions across all episodes. The root cause was `ENTROPY_GATE = 0.8`: the policy skips the ask branch when $H(b) \leq$ `ENTROPY_GATE`, but the uniform prior over K=2 goals has entropy $\ln 2 \approx 0.693 < 0.8$, so the gate fired at every step. After reducing the threshold to 0.3 (below $\ln 2$ but above the near-certain regime), *AskOrAct* achieves 97% success vs. 85% for *NeverAsk* at K=2, ε=0, β=1. Three additional minor bugs were fixed: a NaN guard for unreachable goal positions, a degenerate-episode sentinel, and a warning for unregistered question IDs.

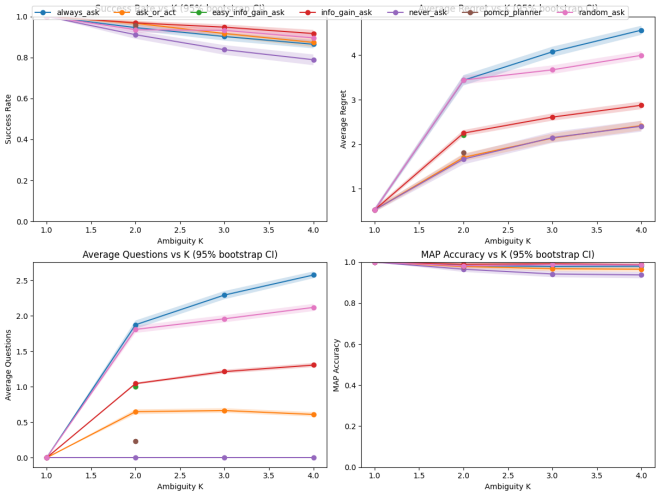

Figure 1: Main evaluation dashboard (100 ep/cell; error bars = 95% CI). Top: success rate. Middle: regret. Bottom: questions asked per episode.

Preliminary Findings

Figure 1 and Tables 1–2 summarise the key results. At K=1 (no ambiguity) all policies are equivalent — a sanity check confirming correct implementation. At $K \geq 3$ instruction ambiguity becomes genuine and policies diverge substantially.

AskOrAct outperforms *NeverAsk* with statistical significance at both high-ambiguity levels: +7.9 pp at K=3 (95% CI $[5.0, 10.8]$) and +8.4 pp at K=4 (95% CI $[5.1, 12.1]$), while asking only 0.6–0.7 questions per episode on average. Both confidence intervals exclude zero, confirming the effect is not due to sampling noise. *AlwaysAsk* asks 3–4$\times$ more per episode but achieves similar or lower success and +2.1 higher regret at K=4, confirming that over-asking is costly.

Table 1: Results at K=3, ε=0, β=1 (100 episodes).

Policy	Success	Avg Q	Regret
AskOrAct	**90%**	0.74	**2.19**
NeverAsk	76%	0.00	2.51
InfoGainAsk	93%	1.25	2.58
AlwaysAsk	87%	2.41	4.14

Table 2: K=2, ε=0, β=1, post entropy-gate fix (100 ep.).

Policy	Success	Avg Q	Regret
AskOrAct	**97%**	0.71	**1.84**
NeverAsk	85%	0.00	2.16
InfoGainAsk	95%	1.04	2.44
AlwaysAsk	91%	1.97	3.67

Robustness (positive and negative findings). Answer-noise sweeps show the *AskOrAct* advantage holds at K=4 across all noise levels (+4–+9 pp), but reverses briefly at K=3, noise = 0.1 (92.0% vs. 94.0% for *NeverAsk*) — noisy answers

can mislead the posterior. Under principal model mismatch, *AskOrAct* retains $+13.5$ pp even when the principal acts near-randomly ($\beta_{\text{true}}=1$ vs. $\hat{\beta}=2$). The ablation confirms a non-zero question cost is essential: $c_q=0$ causes over-asking (1.61 vs. 0.65 q/ep) and reduces success from 73.8% to 69.2%. On the success–regret Pareto frontier at $K=4$, *AskOrAct* is efficient: it dominates *NeverAsk* in success and *AlwaysAsk*/*RandomAsk* in regret.

Key Plot Description for Final Report

Figure 1 is the central result figure. The horizontal axis is ambiguity level $K \in \{1, 2, 3, 4\}$; three stacked panels show (top) success rate, (middle) regret, and (bottom) average questions per episode, with 95% CI error bars. Readers see: (a) all policies tie at $K=1$ (sanity check); (b) the success gap between *AskOrAct* and *NeverAsk* widens with ambiguity; (c) *AlwaysAsk*/*RandomAsk* ask far more questions without a commensurate gain.

An **aggregate delta table** (planned for final report) will have rows $K \in \{2, 3, 4\}$ and columns ΔSuccess (*AskOrAct*$-$*NeverAsk*) with 95% CI, and ΔRegret (*AskOrAct*$-$*AlwaysAsk*) with questions/ep, letting readers see at a glance which improvements are statistically significant.

Plan for Remaining Time

Step 1 — Failure-mode analysis (*by March 21*). Break down timeout vs. wrong-pick failures per (K, ε, β) for all policies. Pinpoint conditions where *AskOrAct* regresses (notably $K=3$, noise $= 0.1$) and determine whether failures stem from posterior misdirection, exhausted question budget, or an overly restrictive ask window.

Step 2 — Final report (*by April 3*). Write the 4-page paper covering introduction, related work, method (posterior update, VoI rule, gating), experiments (aggregate tables, delta-CI analysis), ablations, robustness, and conclusion.

Step 3 — Presentation (*by April 8*). Prepare 14 slides with backup slides for likely Q&A topics (entropy gate sensitivity, question cost ablation, scalability to $K>4$).

Step 4 — Submission package (*by April 10*). Run `./run.sh package` to produce `submission_package.zip` with all CSVs, plots, PDFs, and a `REPRODUCE.txt` with checksums for all seven output files.

AI Tool Declaration

I used Claude Sonnet 4.6 (Anthropic) for coding assistance, debugging, and text editing throughout the project. I am responsible for all research decisions, experimental design, and the content of this report.

References

[1] Chris L. Baker, Rebecca Saxe, and Joshua B. Tenenbaum. Action understanding as inverse planning. *Cognition*, 113 (3):329–349, 2009.

[2] Dylan Hadfield-Menell, Anca D. Dragan, Pieter Abbeel, and Stuart Russell. Cooperative inverse reinforcement learning. In *Advances in Neural Information Processing Systems*, 2016.

[3] Julian Jara-Ettinger, Hyowon Gweon, Laura E. Schulz, and Joshua B. Tenenbaum. The naive utility calculus. *Trends in Cognitive Sciences*, 20(8):589–604, 2016.

[4] Zhi-Xuan Tan, Lexing Ying, Vikash Mansinghka, and Joshua B. Tenenbaum. Pragmatic instruction following and goal assistance via cooperative language guided inverse planning. In *Proceedings of the International Conference on Autonomous Agents and Multiagent Systems*, 2024.